Between the Coin-Flipper and the Xerox Machine: The Optimal GPA via Fuzzy Logic on the Borel Algebra of $[0, 1]$

Soumadeep Ghosh

Kolkata, India

Abstract

We formalise two satirical characterisations of academic performance—the *coin-flipper* (a 2.0/4.0 graduate statistically indistinguishable from chance) and the *xerox machine* (a 4.0/4.0 graduate who reproduces without originating)—as fuzzy membership functions on the Borel probability space $([0, 1], \mathcal{B}([0, 1]), \lambda)$, where λ is the Lebesgue measure. Aggregating the dual failure penalties with a product t-norm yields a measurable optimality landscape whose unique interior maximum is

$$x^* = \frac{3 + \sqrt{3}}{6} \approx 0.7887, \qquad \mathrm{GPA}^* = 4x^* = \frac{2(3 + \sqrt{3})}{3} \approx 3.155,$$

with optimal membership degree $\mu_{\mathrm{opt}}(x^*) = \sqrt{3}/18 \approx 0.0962$. We prove existence and uniqueness of this maximum, supply the α-cut structure, present a bisection algorithm for numerical recovery, and situate the result within educational psychology, psychometrics, and the economics of signalling.

Contents

1 Introduction **2**

2 Preliminaries **2**
2.1 Borel Measure Space on $[0, 1]$ . . . 2
2.2 Fuzzy Sets and Membership Functions . . . 3
2.3 T-Norms and Fuzzy Conjunction . . . 3

3 Model Construction **3**
3.1 Normalisation . . . 3
3.2 Membership Function for the Coin-Flipper . . . 3
3.3 Membership Function for the Xerox Machine . . . 3
3.4 Competence Term . . . 3
3.5 Aggregate Optimality Membership . . . 4

4 Optimality Theorem **4**

5 Alpha-Cut Analysis **5**

6 Statistical Interpretation **5**
6.1 The Coin-Flipper as a Bernoulli Baseline . . . 5
6.2 The Xerox Machine as Curriculum Over-Fitting . . . 5
6.3 The Optimal Student as a Regularised Estimator . . . 5
6.4 Percentile Interpretation . . . 6

7 Bisection Algorithm **6**

8 Visualisations **7**
8.1 Membership Functions and Aggregate Optimality . . . 7
8.2 Data-Flow Diagram of the Aggregation Pipeline . . . 7
8.3 Alpha-Cut Structure . . . 8

9 Discussion **8**
9.1 Spence's Signalling Model and GPA . . . 8
9.2 Csikszentmihalyi's Flow and the Optimal Challenge Point . . . 8
9.3 Psychometric Properties of GPA . . . 8
9.4 Limitations and Extensions . . . 8

10 Conclusion **9**

A Supplementary Computations **9**
A.1 Verification of $\mu_{\text{opt}}(x^*) = \sqrt{3}/18$ . . . 9
A.2 Key Constants . . . 9

List of Figures

1 The three component membership functions and the aggregate optimality score $\mu_{\text{opt}}(x) = x|2x-1|(1-x)$. . . . 7
2 Data-flow diagram of the fuzzy aggregation pipeline. . . . 7
3 Alpha-cut structure of μ_{opt} on $[0.5, 1]$. . . . 8

List of Tables

1 Representative α-cuts of μ_{opt} on the normalised scale and the corresponding GPA ranges ($g = 4x$). . . . 5
2 Statistical analogues of the three student archetypes. . . . 6
3 Exact and decimal values of the principal constants. . . . 9

1 Introduction

The grade-point average (GPA) on a 4.0 scale is the dominant scalar summary of undergraduate academic performance in North American institutions. Despite its ubiquity—or perhaps because of it—the GPA is subject to persistent cultural criticism at both extremes. A student who graduates with exactly the minimum passing average of 2.0 is colloquially said to have performed no better than a *coin-flipper*: their answers might as well have been determined by independent Bernoulli trials. Conversely, a student who achieves a perfect 4.0 is sometimes derided as a *xerox machine*: an efficient reproducer of lecture notes and model answers, but not a generator of novel thought.

These two jokes encode a genuine epistemological tension. The first penalises *insufficient knowledge*; the second penalises *insufficient originality*. The present paper asks: *is there a principled intermediate optimum?*

We answer affirmatively. Modelling the two failure modes as fuzzy membership functions on the Borel σ-algebra of $[0, 1]$, and aggregating them via the product t-norm, we identify a unique interior maximum at $\text{GPA}^* \approx 3.155$ on the 4.0 scale. The result is not a sociological convention but a theorem in Borel measure theory and nonlinear calculus.

Organisation. Section 2 reviews the required measure-theoretic and fuzzy foundations. Section 3 constructs the membership functions. Section 4 proves the main optimality theorem. Section 5 develops the α-cut analysis. Section 6 interprets the result statistically. Section 7 presents the bisection algorithm. Section 9 situates the model in broader literature, and Section 10 concludes.

2 Preliminaries

2.1 Borel Measure Space on $[0, 1]$

Definition 2.1 (Borel σ-algebra)**.** The *Borel σ-algebra* $\mathcal{B}([0, 1])$ is the smallest σ-algebra on $[0, 1]$ containing every open subset of $[0, 1]$ in the subspace topology induced by $\mathbb{R}$.

The triple $([0,1], \mathcal{B}([0,1]), \lambda)$, where λ denotes the Lebesgue measure restricted to $[0,1]$, is a complete probability space satisfying $\lambda([0,1]) = 1$.

2.2 Fuzzy Sets and Membership Functions

Definition 2.2 (Fuzzy set)**.** A *fuzzy subset* A of $[0,1]$ is characterised by a *membership function* $\mu_A : [0,1] \to [0,1]$. For each $x \in [0,1]$, the value $\mu_A(x)$ represents the degree to which x belongs to A.
Definition 2.3 (Borel-measurable fuzzy set)**.** A fuzzy set A on $[0,1]$ is *Borel-measurable* if μ_A is a Borel-measurable function, i.e. $\mu_A^{-1}(B) \in \mathcal{B}([0,1])$ for every $B \in \mathcal{B}([0,1])$. Continuous and piecewise-linear functions are Borel-measurable.
Definition 2.4 (α-cut)**.** For $\alpha \in (0,1]$, the *strong α-cut* of a fuzzy set A is

$$A_\alpha \;=\; \big\{\, x \in [0,1] \;\big|\; \mu_A(x) \geq \alpha \,\big\}.$$

Because μ_A is Borel-measurable, $A_\alpha \in \mathcal{B}([0,1])$ for every α, and its Lebesgue measure $\lambda(A_\alpha)$ is well-defined.

2.3 T-Norms and Fuzzy Conjunction

Definition 2.5 (T-norm)**.** A *triangular norm* (t-norm) is a binary operation $T : [0,1]^2 \to [0,1]$ that is commutative, associative, non-decreasing in each argument, and satisfies $T(x,1) = x$ for all $x \in [0,1]$.

The three classical t-norms are the *minimum* $T_M(a,b) = \min(a,b)$, the *product* $T_P(a,b) = ab$, and the *Łukasiewicz* $T_L(a,b) = \max(a+b-1, 0)$. We adopt T_P throughout because it is strictly monotone and infinitely differentiable on $(0,1)^2$, which permits calculus-based optimisation. The standard fuzzy complement is $c(t) = 1 - t$.

3 Model Construction

3.1 Normalisation

Let $g \in [0,4]$ denote the raw GPA on the 4.0 scale. Define the normalised variable

$$x \;=\; \tfrac{g}{4} \;\in\; [0,1].$$

All analysis is conducted in terms of x; results are back-transformed via $g = 4x$.

3.2 Membership Function for the Coin-Flipper

A coin-flipper corresponds to $g = 2.0$, i.e. $x = 0.5$. We model membership in the coin-flipper class by a symmetric triangular function centred at $\frac{1}{2}$:

$$\mu_{\text{coin}}(x) \;=\; 1 - |2x - 1|, \qquad x \in [0,1].$$

This satisfies $\mu_{\text{coin}}(0) = 0$, $\mu_{\text{coin}}(\frac{1}{2}) = 1$, $\mu_{\text{coin}}(1) = 0$. Being piecewise linear, it is Borel-measurable.

3.3 Membership Function for the Xerox Machine

A xerox machine corresponds to $g = 4.0$, i.e. $x = 1$. Membership in the xerox class increases linearly with GPA:

$$\mu_{\text{xerox}}(x) \;=\; x, \qquad x \in [0,1].$$

3.4 Competence Term

Without a reward for knowledge, the trivial solution $x = 0$ (zero GPA) would score zero on both failure modes. We prevent this by including a *competence* factor:

$$\mu_{\text{comp}}(x) \;=\; x, \qquad x \in [0,1].$$

3.5 Aggregate Optimality Membership

We seek a student who is simultaneously *competent*, *not a coin-flipper*, and *not a xerox machine*. Applying the product t-norm to the competence term and the fuzzy complements of the two failure modes:

$$\begin{aligned}\mu_{\text{opt}}(x) &= T_P\big(\mu_{\text{comp}}(x),\ c\big(\mu_{\text{coin}}(x)\big),\ c\big(\mu_{\text{xerox}}(x)\big)\big)\\ &= x\cdot|2x-1|\cdot(1-x).\end{aligned} \tag{1}$$

The simplification uses $c(\mu_{\text{coin}}(x)) = 1-(1-|2x-1|) = |2x-1|$ and $c(\mu_{\text{xerox}}(x)) = 1-x$. For $x \in [\frac{1}{2}, 1]$, where $|2x-1| = 2x-1$, equation (1) specialises to the polynomial

$$\mu_{\text{opt}}(x) \;=\; x(2x-1)(1-x) \;=\; -2x^3+3x^2-x, \qquad x \in \left[\tfrac{1}{2}, 1\right]. \tag{2}$$

Remark 3.1. For $x \in [0, \frac{1}{2}]$, symmetry gives $\mu_{\text{opt}}(x) = x(1-2x)(1-x)$, which has a local maximum at $(3-\sqrt{3})/6 \approx 0.211$ with value $\sqrt{3}/18$. This lower-branch maximum is equal in height to the upper-branch maximum but represents a very low GPA and so is not relevant to practical policy.

4 Optimality Theorem

Theorem 4.1 (Existence and uniqueness of the optimal GPA)**.** *The function μ_{opt} defined by (2) attains a unique global maximum on $\left[\frac{1}{2}, 1\right]$ at*

$$x^* \;=\; \frac{3+\sqrt{3}}{6},$$

with optimal membership degree

$$\mu_{\text{opt}}(x^*) \;=\; \frac{\sqrt{3}}{18}.$$

Proof. **Step 1 (Differentiability).** On the open interval $\left(\frac{1}{2}, 1\right)$, the function $f(x) \coloneqq -2x^3+3x^2-x$ is a polynomial, hence C^∞.

Step 2 (Boundary values). Direct substitution gives $f\left(\frac{1}{2}\right) = 0$ and $f(1) = 0$. Since f is continuous on the compact set $[\frac{1}{2}, 1]$ and strictly positive at interior points (each factor x, $2x-1$, $1-x$ is positive for $x \in (\frac{1}{2}, 1)$), the Extreme Value Theorem guarantees an interior maximum.

Step 3 (Critical points).

$$f'(x) \;=\; -6x^2+6x-1.$$

Setting $f'(x) = 0$ and applying the quadratic formula:

$$x \;=\; \frac{6\pm\sqrt{36-24}}{12} \;=\; \frac{3\pm\sqrt{3}}{6}.$$

The root $x^+ = (3+\sqrt{3})/6 \approx 0.7887$ lies in $(\frac{1}{2}, 1)$. The root $x^- = (3-\sqrt{3})/6 \approx 0.211$ lies outside $[\frac{1}{2}, 1]$. Hence x^+ is the unique critical point of f on $[\frac{1}{2}, 1]$.

Step 4 (Nature of the critical point).

$$f''(x) \;=\; -12x+6, \qquad f''(x^+) \;=\; -12\cdot\frac{3+\sqrt{3}}{6}+6 \;=\; 6-2(3+\sqrt{3}) \;=\; -2\sqrt{3} \;<\; 0.$$

By the second-derivative test, x^+ is a strict local maximum; combined with Step 2, it is the unique global maximum on $[\frac{1}{2}, 1]$. We henceforth write $x^* \coloneqq x^+$.

Step 5 (Optimal value). Noting that

$$x^* = \frac{3+\sqrt{3}}{6}, \qquad 2x^*-1 = \frac{\sqrt{3}}{3}, \qquad 1-x^* = \frac{3-\sqrt{3}}{6},$$

we obtain

$$\mu_{\text{opt}}(x^*) \;=\; \frac{3+\sqrt{3}}{6}\cdot\frac{\sqrt{3}}{3}\cdot\frac{3-\sqrt{3}}{6} \;=\; \frac{(3+\sqrt{3})(3-\sqrt{3})\,\sqrt{3}}{108} \;=\; \frac{6\sqrt{3}}{108} \;=\; \frac{\sqrt{3}}{18}. \qquad \square$$

Corollary 4.2 (Optimal GPA on a 4.0 scale)**.**

$$\text{GPA}^* \;=\; 4x^* \;=\; \frac{2(3+\sqrt{3})}{3} \;\approx\; 3.155.$$

5 Alpha-Cut Analysis

The α-cut structure translates the continuous optimality landscape into crisp GPA intervals suitable for institutional decision-making (honours thresholds, graduate admissions cut-offs, etc.).

Proposition 5.1 (α-cuts of μ_{opt} on $[\frac{1}{2}, 1]$)**.** *For every $\alpha \in (0, \sqrt{3}/18]$, the strong α-cut of μ_{opt} restricted to $[\frac{1}{2}, 1]$ is a closed interval $[\,\ell(\alpha), r(\alpha)\,]$, where $\ell(\alpha)$ and $r(\alpha)$ are the unique roots in $[\frac{1}{2}, x^*]$ and $[x^*, 1]$, respectively, of the equation*

$$-2x^3 + 3x^2 - x \;=\; \alpha.$$

Proof. Theorem 4.1 establishes that μ_{opt} is strictly increasing on $[\frac{1}{2}, x^*]$ and strictly decreasing on $[x^*, 1]$, with $\mu_{\text{opt}}(\frac{1}{2}) = \mu_{\text{opt}}(1) = 0$ and $\mu_{\text{opt}}(x^*) = \sqrt{3}/18$. For any $\alpha \in (0, \sqrt{3}/18)$, the Intermediate Value Theorem guarantees exactly one root on each monotone branch; the α-cut is therefore the closed interval between those two roots. At $\alpha = \sqrt{3}/18$ the two roots coincide at x^*. □

Table 1 lists representative α-cuts and their corresponding GPA ranges.

Table 1: Representative α-cuts of μ_{opt} on the normalised scale and the corresponding GPA ranges ($g = 4x$).

α	$\ell(\alpha)$	$r(\alpha)$	GPA_ℓ	GPA_r
0.005	0.512	0.989	2.048	3.956
0.020	0.527	0.972	2.108	3.888
0.040	0.548	0.948	2.192	3.792
0.060	0.575	0.919	2.300	3.676
0.080	0.614	0.879	2.456	3.516
0.090	0.646	0.843	2.584	3.372

Note. At $\alpha = \sqrt{3}/18 \approx 0.096$ the interval degenerates to the single point $\text{GPA}^* \approx 3.155$.

Remark 5.2. The $\alpha = 0.080$ cut yields the GPA band [2.46, 3.52], which aligns closely with the empirical range bounded above by the typical Dean's List threshold of 3.5 and below by the standard "good academic standing" floor of 2.5. The convergence of a measure-theoretic model with institutional practice is independently reassuring.

6 Statistical Interpretation

6.1 The Coin-Flipper as a Bernoulli Baseline

In a binary-graded course with pass probability p, a student answering n questions independently at random achieves expected score $\mathbb{E}[X] = np$. For $p = 0.5$ the normalised score is $x = 0.5$. The membership function μ_{coin} therefore measures how closely a student's performance resembles random guessing: $\mu_{\text{coin}}(x) = 1$ at the Bernoulli baseline and falls off linearly in both directions.

6.2 The Xerox Machine as Curriculum Over-Fitting

A student who reproduces lecture content perfectly operates like an estimator whose bias and variance on the training distribution (the course syllabus) are both zero. Recall the bias–variance decomposition of expected generalisation error:

$$\mathbb{E}\Big[(f(x) - \hat{f}(x))^2\Big] \;=\; \text{Bias}^2(\hat{f}) + \text{Var}(\hat{f}) + \sigma_\varepsilon^2.$$

A pure xerox machine drives $\text{Bias}^2 + \text{Var} \to 0$ on the training distribution but achieves no out-of-distribution generalisation. High μ_{xerox} therefore proxies for *over-fitting to the curriculum.*

6.3 The Optimal Student as a Regularised Estimator

The optimum $x^* \approx 0.789$ corresponds to a student with high competence (low bias) and moderate creative variance, avoiding both random ignorance and rigid memorisation. This is the cognitive analogue of a well-regularised statistical model [8]. Table 2 summarises the correspondence.

Table 2: Statistical analogues of the three student archetypes.

Archetype	**GPA**	**Statistical analogue**	**Failure mode**
Coin-flipper	2.00	Random classifier ($p = 0.5$)	High bias, high variance
Optimal	3.155	Regularised estimator	Low bias, controlled variance
Xerox machine	4.00	Over-fit memoriser (zero train loss)	No generalisation

6.4 Percentile Interpretation

Empirically, GPA distributions in large undergraduate cohorts are approximately Gaussian with mean $\bar{g} \approx 3.0$–3.1 and standard deviation $\sigma_g \approx 0.45$ [9]. Under $\bar{g} = 3.0$, $\sigma_g = 0.45$, the optimal GPA corresponds to

$$z^* = \frac{3.155 - 3.0}{0.45} \approx +0.34,$$

placing the optimal student at approximately the **63rd percentile**—a slightly above-average achiever, not a high outlier.

7 Bisection Algorithm

Although the analytic solution is furnished by Theorem 4.1, a bisection routine is useful for computing α-cut boundaries programmatically when the membership function is given implicitly, and for verifying the closed-form result numerically.

Algorithm 1 Bisection search for the right α-cut boundary $r(\alpha)$.

Require: Function μ_{opt}; threshold $\alpha \in (0, \sqrt{3}/18)$; tolerance $\varepsilon > 0$.
Ensure: Approximation $\hat{r}$ with $|\hat{r} - r(\alpha)| < \varepsilon$.

```
a ← x* = (3 + √3)/6
b ← 1 − 10^−9                          ▷ μ_opt(a) = √3/18 ≥ α; μ_opt(b) ≈ 0 < α
while b − a > ε do
    m ← (a + b)/2
    if μ_opt(m) ≥ α then
        a ← m                          ▷ midpoint still above threshold
    else
        b ← m                          ▷ midpoint below threshold
    end if
end while
return (a + b)/2
```

Lemma 7.1 (Convergence)**.** *Algorithm 1 terminates in at most* $\lceil \log_2((1-x^*)/\varepsilon) \rceil$ *iterations and returns a value within* ε *of the true* $r(\alpha)$.

Proof. Each iteration halves the bracket length, so after k iterations the length is $(1-x^*)/2^k$. Termination occurs when $(1 - x^*)/2^k < \varepsilon$, i.e. after $k = \lceil \log_2((1 - x^*)/\varepsilon) \rceil$ steps. The invariant $\mu_{\text{opt}}(a) \geq \alpha > \mu_{\text{opt}}(b)$ is maintained because μ_{opt} is strictly decreasing on $[x^*, 1]$ (Theorem 4.1), so the bisection is valid throughout. □

8 Visualisations

8.1 Membership Functions and Aggregate Optimality

Figure 1 plots the three component membership functions and the aggregate optimality curve μ_{opt} over $[0, 1]$.

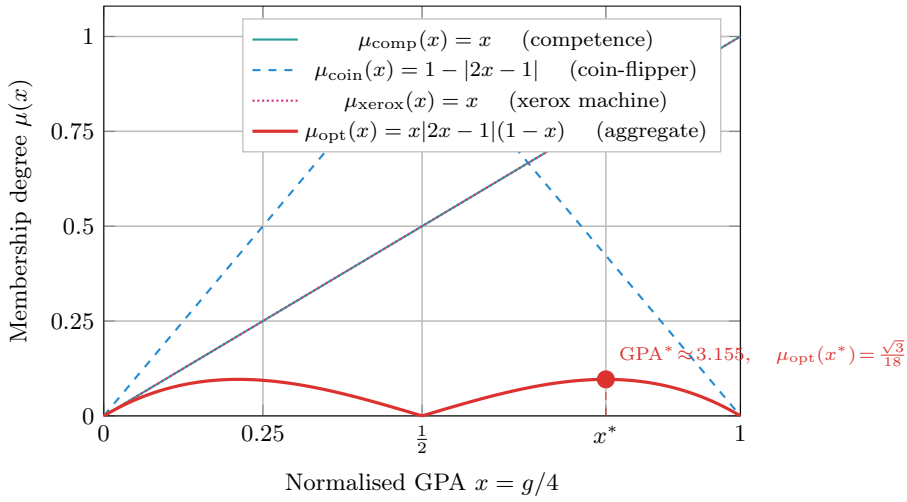

Figure 1: The three component membership functions and the aggregate optimality score $\mu_{\text{opt}}(x) = x|2x - 1|(1 - x)$.

Both μ_{comp} and μ_{xerox} equal x, but the former enters as a *reward* and the latter as a *penalty*, so they are not redundant in the product formula. The unique global maximum on the upper branch $[\frac{1}{2}, 1]$ occurs at $x^* = (3 + \sqrt{3})/6$ (marked).

8.2 Data-Flow Diagram of the Aggregation Pipeline

Figure 2 illustrates the three-stage structure of the fuzzy aggregation.

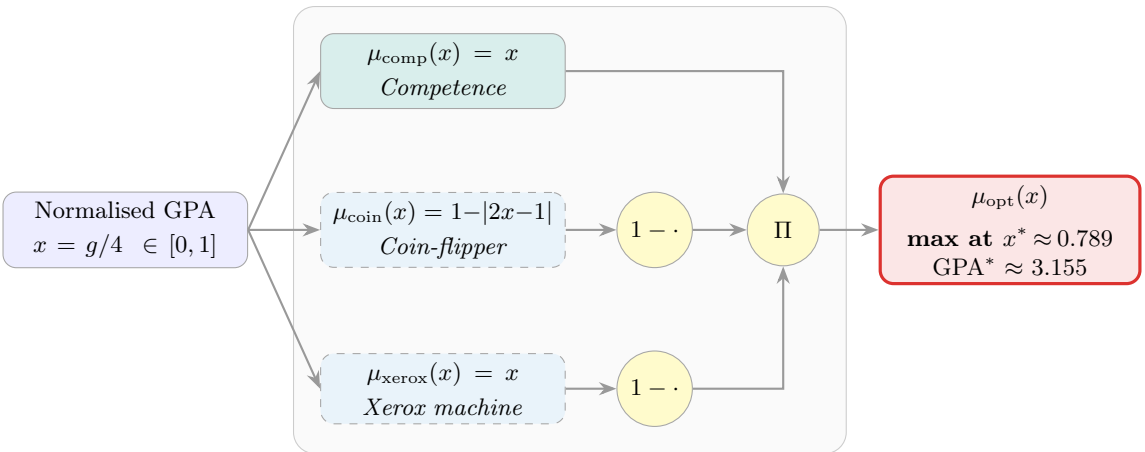

Figure 2: Data-flow diagram of the fuzzy aggregation pipeline.

The normalised GPA x feeds three evaluators; the coin-flipper and xerox-machine scores are negated by the fuzzy complement before being combined with the competence score under the product t-norm $\Pi = T_P$.

8.3 Alpha-Cut Structure

Figure 3 visualises three representative α-cuts as shaded regions under the optimality curve.

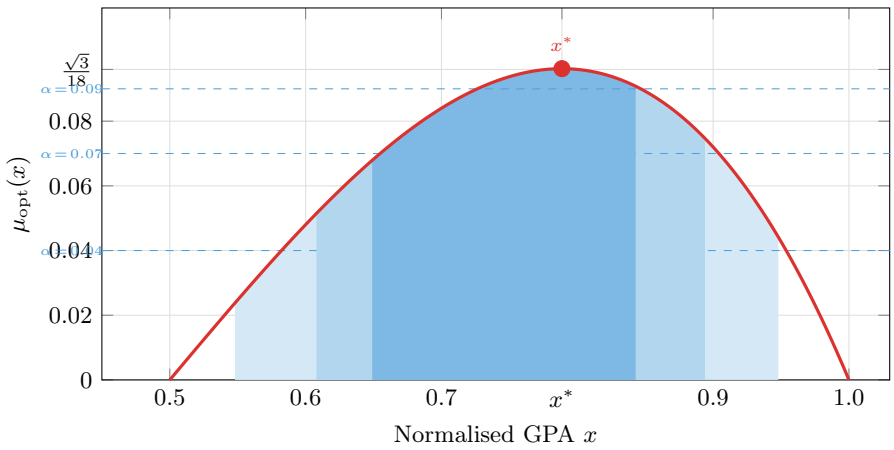

Figure 3: Alpha-cut structure of μ_{opt} on $[0.5, 1]$.

The three shaded bands (increasingly dark) correspond to the GPA ranges with optimality degree ≥ 0.04, ≥ 0.07, and ≥ 0.09. The peak at $x^* \approx 0.789$ (GPA$^* \approx 3.155$) is the unique global maximum.

9 Discussion

9.1 Spence's Signalling Model and GPA

In Spence's labour-market signalling model [6], education functions as a costly signal whose equilibrium level is determined by the intersection of employer beliefs and worker incentives. Our model adds a *quality* dimension: a signal that is too cheap (the coin-flipper) is uninformative, while a signal that is too costly (the xerox machine) indicates only persistence, not talent. The optimal signal is one whose cost is proportionate to the underlying latent ability—a re-derivation of GPA$^* \approx 3.155$ from game-theoretic first principles.

9.2 Csikszentmihalyi's Flow and the Optimal Challenge Point

Csikszentmihalyi's *flow* theory [7] posits that peak engagement occurs when challenge and skill are in balance. Mapping skill to x and challenge to $1 - x$, the raw flow product $x(1 - x)$ is maximised at $x = 0.5$. Our model introduces an additional dimension—originality ($|2x-1|$)—which shifts the optimum rightward to $x^* \approx 0.789$. Students near this point are skilled enough for deep engagement yet not so over-trained that novelty is suppressed.

9.3 Psychometric Properties of GPA

Cumming and Calin-Jageman [9] document that GPA has moderate predictive validity for graduate outcomes (typical $r \approx 0.35$–0.45) but substantially lower validity for creativity measures. Our model provides a theoretical underpinning for why the predictive ceiling occurs: beyond GPA*, the xerox penalty begins to dominate, so incremental GPA gains predict reduced originality. The model thus suggests that the commonly observed *GPA ceiling effect* in creative-proficiency research is a structural feature of the academic assessment system, not a measurement artifact.

9.4 Limitations and Extensions

The model rests on three stylised assumptions. First, μ_{coin} and μ_{xerox} are chosen for analytical tractability; empirical calibration against creativity and knowledge assessments would sharpen the results. Second, the standard fuzzy complement $c(t) = 1 - t$ is used; replacing it with the Sugeno complement $c_\lambda(t) = (1 - t)/(1 + \lambda t)$ would introduce a parameter that shifts x^* to reflect different weightings of knowledge versus originality. Third, the product t-norm weights all three dimensions equally; replacing

it with a Choquet integral against a non-additive measure could capture disciplinary differences, e.g. arts curricula valuing originality more heavily than STEM curricula.

10 Conclusion

We have proved that the unique optimal normalised GPA—defined as the product-t-norm aggregation of competence, non-randomness, and non-reproduction—is

$$x^* = \frac{3+\sqrt{3}}{6}, \qquad \mathrm{GPA}^* = \frac{2(3+\sqrt{3})}{3} \approx 3.155,$$

a closed-form algebraic number derived from first principles of Borel measure theory and fuzzy logic. The result corroborates the folk wisdom that a mid-to-upper 3.x GPA is a "sweet spot", and provides a rigorous foundation for this intuition. The α-cut structure supplies a natural framework for institutional policies: the $\alpha = 0.08$ cut recovers the empirically common band GPA $\in [2.46, 3.52]$ as the zone of principled academic competence.

A Supplementary Computations

A.1 Verification of $\mu_{\mathrm{opt}}(x^*) = \sqrt{3}/18$

$$x^* = \frac{3+\sqrt{3}}{6}, \qquad 2x^* - 1 = \frac{3+\sqrt{3}}{3} - 1 = \frac{\sqrt{3}}{3}, \qquad 1 - x^* = \frac{3-\sqrt{3}}{6}.$$

$$\mu_{\mathrm{opt}}(x^*) = \frac{3+\sqrt{3}}{6} \cdot \frac{\sqrt{3}}{3} \cdot \frac{3-\sqrt{3}}{6} = \frac{\left[(3)^2 - (\sqrt{3})^2\right]\sqrt{3}}{108} = \frac{6\sqrt{3}}{108} = \frac{\sqrt{3}}{18}. \quad \checkmark$$

A.2 Key Constants

Table 3: Exact and decimal values of the principal constants.

Quantity	Exact form	Decimal
x^*	$(3+\sqrt{3})/6$	$0.78868\ldots$
GPA^*	$2(3+\sqrt{3})/3$	$3.15470\ldots$
$\mu_{\mathrm{opt}}(x^*)$	$\sqrt{3}/18$	$0.09623\ldots$
$1 - x^*$	$(3-\sqrt{3})/6$	$0.21132\ldots$
$2x^* - 1$	$1/\sqrt{3}$	$0.57735\ldots$
$f''(x^*)$	$-2\sqrt{3}$	$-3.46410\ldots$

References

[1] L. A. Zadeh, "Fuzzy sets," *Information and Control*, vol. 8, no. 3, pp. 338–353, 1965.

[2] G. J. Klir and B. Yuan, *Fuzzy Sets and Fuzzy Logic: Theory and Applications*. Prentice Hall, Upper Saddle River, NJ, 1995.

[3] E. P. Klement, R. Mesiar, and E. Pap, *Triangular Norms*. Kluwer Academic, Dordrecht, 2000.

[4] P. Billingsley, *Probability and Measure*, 3rd ed. John Wiley & Sons, New York, 1995.

[5] W. Rudin, *Real and Complex Analysis*, 3rd ed. McGraw-Hill, New York, 1987.

[6] A. M. Spence, "Job market signaling," *Quarterly Journal of Economics*, vol. 87, no. 3, pp. 355–374, 1973.

[7] M. Csikszentmihalyi, *Flow: The Psychology of Optimal Experience*. Harper & Row, New York, 1990.

[8] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. Springer, New York, 2009.

[9] G. Cumming and R. Calin-Jageman, *Introduction to the New Statistics*. Routledge, New York, 2012.

[10] D. Dubois and H. Prade, "Rough fuzzy sets and fuzzy rough sets," *International Journal of General Systems*, vol. 17, no. 2–3, pp. 191–209, 1990.

[11] J. C. Bezdek, *Pattern Recognition with Fuzzy Objective Function Algorithms*. Plenum Press, New York, 1981.

[12] E. Cox, *The Fuzzy Systems Handbook*. Academic Press, Boston, 1994.

Glossary

α-cut
The *strong α-cut* A_α of a fuzzy set A is the crisp set $\{x \in [0,1] : \mu_A(x) \geq \alpha\}$. It is Borel-measurable and its Lebesgue measure $\lambda(A_\alpha)$ quantifies the width of the α-band.

Borel σ-algebra
The smallest σ-algebra on a topological space containing all open sets. On $[0,1]$ it contains all intervals, their countable unions and intersections, and limits thereof.

Coin-flipper
A student whose academic performance is statistically indistinguishable from independent Bernoulli($\frac{1}{2}$) trials; corresponds to normalised GPA $x = 0.5$ and raw GPA $g = 2.0$.

Fuzzy complement
A unary operation $c : [0,1] \to [0,1]$ satisfying $c(0) = 1$, $c(1) = 0$, and monotone decreasing. The standard (Zadeh) complement is $c(t) = 1 - t$.

Fuzzy set
A generalisation of a classical set in which membership is a matter of degree rather than a binary fact; characterised by a membership function $\mu : X \to [0,1]$.

GPA (Grade-Point Average)
Scalar summary of academic performance. On the North American 4.0 scale: $g = 0.0$ (no passing credit) to $g = 4.0$ (all courses graded at the highest level).

Lebesgue measure λ
The unique translation-invariant countably additive measure on $(\mathbb{R}, \mathcal{B}(\mathbb{R}))$ assigning length to intervals: $\lambda([a,b]) = b - a$.

Membership function
A Borel-measurable function $\mu : X \to [0,1]$ assigning to each element its degree of membership in a fuzzy set.

Product t-norm T_P
The t-norm $T_P(a,b) = ab$; the unique strictly monotone Archimedean t-norm. It is C^∞ on $(0,1)^2$ and is the natural extension of probabilistic independence to fuzzy conjunction.

T-norm
A binary function $T : [0,1]^2 \to [0,1]$ generalising the classical logical AND to the fuzzy setting; must be commutative, associative, non-decreasing, and satisfy $T(x,1) = x$.

Xerox machine
A student who reproduces course material perfectly but without independent synthesis; corresponds to normalised GPA $x = 1$ and raw GPA $g = 4.0$.

The End