ABUNDANCE BY DESIGN

What Happens When AI Optimizes for Human Flourishing

A Quantitative, Verbal, and Critical Analysis

"The test of our progress is not whether we add more to the abundance of those who have much; it is whether we provide enough for those who have too little."

— Franklin D. Roosevelt, Second Inaugural Address, 1937

700M+	$15.7T	40%
people in extreme poverty living on under $2.15/day	projected AI boost to global GDP by 2030	crop-yield improvement via AI precision agriculture
World Bank (2024)	*PwC "Sizing the Prize" (2017)*	*In documented pilot programmes*

Kolkata, India

Abstract

This paper examines the potential of artificial intelligence to serve as a structural instrument for poverty eradication, employing quantitative, narrative, and critical modes of analysis in tandem. Drawing on documented pilot programmes across agriculture, healthcare, and financial inclusion, we find that AI-driven interventions have produced measurable improvements in crop yields (30–40%), disease detection rates (40%), and credit access for previously unbankable populations—with over 1.5 million first-time borrowers reached by a single machine learning lending platform in Kenya alone. We situate these gains against a global backdrop in which 700 million people remain in extreme poverty and established development instruments are yielding diminishing returns. The paper then advances a sustained critical analysis of four structural risks inherent to AI-driven development: the definitional instability of "flourishing" as an optimisation target; the extractive dynamics of data collection from Global South communities (data colonialism); the automation paradox, by which AI-driven productivity may eliminate the industrialisation pathway that lifted previous generations out of poverty; and the governance gap in weak-institution contexts where algorithmic opacity has no accountability counterpart. From this tension between demonstrated impact and structural risk, we derive five design principles—proximity of co-creation, data sovereignty, decision legibility, deliberate redistribution of productivity gains, and epistemic humility about the limits of optimisation—that together constitute a framework for deploying AI in service of genuine, equitable human flourishing. We conclude that the eradication of poverty through AI is neither inevitable nor impossible: it is a design choice, and one that requires explicit values, not merely technical capability.

The paper ends with “The End”

Contents

1 The Question We Have Never Been Able to Answer **2**

2 The Numbers: What AI Is Already Doing **2**

2.1 Agriculture and Food Security . . . 2

2.2 Healthcare Access and Disease Burden . . . 3

2.3 Financial Inclusion and Economic Mobility . . . 3

3 The Human Story: What the Data Cannot Capture **4**

3.1 A Farm in Northern Ghana . . . 4

3.2 The Last-Mile Problem, Reimagined . . . 4

4 The Critique: What Optimism Obscures **5**

4.1 Who Defines Flourishing? . . . 5

4.2 The Data-Colonialism Problem . . . 6

4.3 The Automation Paradox . . . 6

4.4 The Governance Gap . . . 7

5 Toward Abundance by Design: A Framework **7**

6 Conclusion: Designing the Abundance We Choose **8**

Sources and Key References **9**

1 The Question We Have Never Been Able to Answer

For most of recorded history, poverty was regarded as the natural condition of humanity—an immutable feature of the human landscape, as reliable as drought or epidemic disease. By the mid-twentieth century, that fatalism had begun to erode. Development economists mapped the mechanics of deprivation. International institutions organised resources. Governments made ambitious promises. Yet at the turn of the twenty-first century, more than one billion people still lived on less than a dollar a day.

Progress since then has been real and worth acknowledging. The share of the global population living in extreme poverty fell from approximately 36 per cent in 1990 to roughly 9 per cent today—a historic reduction by any measure. But that momentum has slowed. Much of the gain was driven by structural economic transformation in China and South-East Asia: a one-time convergence event unlikely to replicate at comparable speed elsewhere. The roughly 700 million people who remain in extreme poverty are now concentrated in places where the established tools—foreign aid, microfinance, trade liberalisation, conditional cash transfers—are yielding diminishing returns.

It is into this landscape that artificial intelligence arrives. Not as a magic wand, but as a category of instrument unlike any development practitioners have previously had at their disposal: one capable of processing millions of data points simultaneously, optimising across systems of staggering complexity, making individualised decisions at scale, and improving through its own outputs. The question is no longer whether AI will touch poverty—it already does. The question is whether it will be *designed* to eradicate it.

Poverty is not a shortage of charity. It is a shortage of optimised systems—and AI is the most powerful optimiser ever built.

2 The Numbers: What AI Is Already Doing

2.1 Agriculture and Food Security

Food insecurity remains the most immediate face of poverty. Approximately 733 million people face hunger daily, with sub-Saharan Africa and South Asia bearing the heaviest burden. The underlying problem is not global caloric supply—Earth currently produces enough food to feed ten billion people—but endemic inefficiency: in distribution logistics, production planning, and most critically, at the level of the individual farm.

AI is addressing each of these bottlenecks with measurable results. Pilot programmes employing AI-driven precision agriculture in Kenya, India, and Bangladesh have documented crop-yield improvements of 30–40 per cent through optimised planting schedules, real-time soil analysis, and predictive irrigation management. Microsoft's AI for Agriculture initiative, deployed across 175,000 farms in India, produced an average income increase of $180 per farmer per season—modest in absolute terms, but representing a 25–30 per cent improvement in annual household income for subsistence cultivators who had never before had access to agronomic advisory services.

At the supply-chain level, AI-driven logistics optimisation has reduced food waste by up to 35 per cent in documented pilots. Approximately one-third of all food produced

globally—some 1.3 billion tonnes annually—is lost or wasted before it reaches a consumer. A sustained 20 per cent reduction in that loss would free enough calories to eliminate hunger across sub-Saharan Africa twice over. The arithmetic is not complicated; the implementation has been.

2.2 Healthcare Access and Disease Burden

The relationship between poverty and poor health is circular and devastating. Poverty limits access to nutrition, clean water, and medical care, thereby causing disease; disease deepens poverty by destroying productivity and depleting household savings. Breaking this cycle requires reaching patients who are geographically remote, economically marginalised, or systematically excluded from formal health systems—precisely the populations least likely to be served by conventional health infrastructure.

AI diagnostics are demonstrating genuine reach. Google's Lymph Node Assistant (LYNA) detects metastatic breast cancer in biopsy slides with accuracy that matches or exceeds specialist pathologists, and has been piloted in resource-constrained settings in sub-Saharan Africa where such specialists are unavailable. In Rwanda and Kenya, AI-powered tuberculosis screening tools have increased detection rates by 40 per cent in communities previously served only by community health workers with limited diagnostic instruments.

The economic consequences of improved detection are not trivial. The World Health Organization estimates that vaccine-preventable diseases alone cost low-income countries $34 billion annually in lost productivity. AI systems that optimise vaccine cold-chain logistics, predict outbreak vectors, and dynamically route mobile health workers have reduced immunisation delivery costs by up to 25 per cent in West African pilots—improvements that translate directly into lives saved and economic capacity preserved.

2.3 Financial Inclusion and Economic Mobility

Approximately 1.4 billion adults worldwide remain unbanked, without access to the basic financial instruments that citizens of wealthy countries take for granted. The barrier is not principally geographical or infrastructural; it is informational. Traditional lending institutions require credit histories, physical collateral, and formal documentation that the global poor cannot provide. The result is systematic exclusion from the formal economy: no savings accounts, no insurance, no access to credit for productive investment.

Machine learning is restructuring the credit equation. Alternative credit-scoring systems—which model mobile-phone usage patterns, transaction histories, social network data, and psychometric assessments—have enabled fintech companies in Kenya, Nigeria, and the Philippines to extend credit to previously unbankable populations. Kenya's M-Shwari platform, powered by machine learning credit models, had disbursed over 1.5 million loans to first-time borrowers by 2023, with default rates remaining below 3 per cent,

competitive with conventional bank lending. The implication is striking: the "unbankable" classification was never primarily about risk. It was about the absence of data infrastructure capable of seeing that risk accurately.

The multiplier effects of financial inclusion are well documented. Access to basic savings accounts increases household investment in education by 15–20 per cent. Access to credit increases small-business formation rates by approximately 30 per cent. Insurance products reduce the probability that a single health shock or weather event will push a household below the poverty line. AI is not merely accelerating transactions; it is constructing the scaffolding of economic resilience.

3 The Human Story: What the Data Cannot Capture

3.1 A Farm in Northern Ghana

Amara Kone grows yams and cassava on a two-hectare plot in Ghana's Upper West Region. Until 2022, she farmed as her mother had, and her mother's mother before her: planting by seasonal intuition, managing pests with whatever inputs the local cooperative could supply, and hoping the rains would arrive on schedule.

In 2022, her cooperative joined a pilot programme deploying a mobile AI advisory tool developed by the agricultural technology nonprofit Farmerline. Each morning, Amara receives a voice message in Dagbani—her first language—specifying when to irrigate, which rows exhibit early signs of fungal stress, and what soil-moisture sensors suggest about that week's water requirements. The system also provides forward-looking yield estimates three months in advance, allowing her to negotiate sale prices with buyers who had previously exploited informational asymmetry to lowball her at the point of harvest.

In 2023, Amara's yield increased by 38 per cent. More consequentially, she sold her crop at prices 22 per cent higher than the year before, because she possessed—before the buyer arrived—regional harvest data indicating relative market scarcity. She used the additional income to pay secondary school fees for her eldest daughter. That daughter is now studying agronomy at a regional university.

This is a story about yams. It is also a story about compounding: the way small efficiency gains, repeated across seasons and across generations, cascade into transformative life trajectories. Amara's story cannot be captured in an aggregate statistic. But it is precisely what aggregate statistics are built from.

3.2 The Last-Mile Problem, Reimagined

Development practitioners have long spoken of the "last mile" problem: the final, most expensive leg of delivering essential services to the most remote and marginalised populations. Roads terminate. Electrical grids reach their limits. Mobile networks grow thin.

The people who most need services are precisely those who are hardest and most costly to reach by conventional means.

AI is not solving the last-mile problem by building more roads. It is solving it by making roads progressively less necessary. Drone-delivery networks guided by AI logistics models now deliver medical supplies, blood products, and vaccines to remote communities in Rwanda, Ghana, and Tanzania within thirty minutes—journeys that previously required six-hour round trips over unpaved terrain. In regions where roads wash out seasonally during monsoon rains, drone delivery is not a novelty; it is the only reliable supply chain that exists.

What makes this possible is not the drone hardware, which has existed for decades. It is the optimisation layer: routing algorithms that account simultaneously for weather conditions, battery charge, cargo weight, and dynamic demand signals, determining not merely how to deliver but what to deliver, when, and in what quantities. The optimisation runs invisibly and continuously, improving with every completed mission.

The most powerful thing AI can do for the poor is not replace human judgement—it is make the quality of human judgement available to people who have never had access to it.

4 The Critique: What Optimism Obscures

4.1 Who Defines Flourishing?

The phrase "optimising for human flourishing" conceals an enormous philosophical and political question: whose definition of flourishing? Every optimisation requires an objective function—a specific, quantifiable target. But human wellbeing is multidimensional, contested, and culturally particular. Income per capita captures one dimension. Life expectancy captures another. Neither captures dignity, autonomy, cultural continuity, or the capacity to live in accordance with one's own conception of a good life.

When AI systems are deployed to "reduce poverty," they typically optimise for proxies—caloric intake, income levels, school enrolment rates. These proxies are not wrong, but they are partial, and partial optimisation can improve the metric while degrading the underlying reality the metric is intended to represent. A community whose incomes rise because its residents now work in a distant factory, separated from their families and their land, may score better on poverty indices while experiencing a genuine deterioration in collective wellbeing.

This is not a hypothetical concern. Development history contains many interventions that succeeded by their own measures and failed by every other. The Green Revolution of the 1960s and 1970s dramatically increased crop yields across South and South-East Asia—and also concentrated land ownership, marginalised smallholder farmers, and displaced

millions of agricultural labourers. AI-powered agricultural optimisation risks replicating these dynamics at greater speed and scale, unless the design process explicitly incorporates the preferences and priorities of the communities being served.

4.2 The Data-Colonialism Problem

Every AI system learns from data. In the context of poverty alleviation, the effectiveness of AI tools depends entirely on the quality and representativeness of training data. Here lies a structural problem the development-AI community has not adequately confronted: the data landscape of the Global South is thin, patchy, systematically biased, and largely owned by institutions headquartered in the Global North.

Health AI systems trained predominantly on data from North American and European populations produce less accurate diagnoses for patients with different genetic profiles, different disease presentations, and different population-level baselines. Credit-scoring models trained on formal transaction records systematically disadvantage people whose economic activity is informal, communal, and cash-based. Agricultural AI tools developed for temperate climates may perform poorly—or dangerously—in the highly variable microclimates of tropical smallholder farming regions.

The deeper concern is proprietary. When AI tools are developed by private companies or foreign non-governmental organisations, data collected from poor communities routinely flows to servers in San Francisco, London, or Zürich, where it trains models sold back to those communities or generates insights monetised by entities that bear no accountability to the people whose lives produced the data. This arrangement—sometimes termed *data colonialism*—replicates historical patterns of resource extraction in digital form.

4.3 The Automation Paradox

Perhaps the sharpest critique of AI-driven poverty reduction concerns employment. The dominant pathway out of poverty in the twentieth century was industrialisation: the migration of workers from subsistence agriculture into low-skill manufacturing, creating income, building transferable skills, and generating the urban middle classes that drove consumer economies across Asia, Europe, and the Americas. That pathway is closing.

Automation has already eliminated large categories of low-skill manufacturing employment in wealthy countries. More troublingly, it is eliminating them in developing countries *before* those countries have completed an industrialisation cycle. Ethiopia, Bangladesh, and Vietnam built nascent garment industries over decades, only to find them threatened by robotic textile machinery. The traditional ladder out of poverty—from farm to factory to office—is losing rungs faster than the people at the bottom can climb.

The McKinsey Global Institute estimates that by 2030, up to 375 million workers globally may need to change occupational categories as a result of automation, with workers in

lower-income countries disproportionately affected because their current employment is concentrated in the routine, repetitive tasks most susceptible to algorithmic substitution. AI that optimises productivity in global supply chains does not automatically distribute those productivity gains to the workers it displaces.

Proponents of AI-driven development correctly observe that automation has historically created more jobs than it destroys: that new industries, unimaginable in advance, eventually absorb the workers that obsolete industries release. This argument is historically defensible but temporally incomplete. Transitions take decades. A worker made redundant at forty-five does not smoothly retrain for software development. A community built around a single industry does not organically diversify. The optimistic long-run view offers cold comfort to the generation that lives through the transition.

4.4 The Governance Gap

AI systems make consequential decisions—about creditworthiness, medical treatment, agricultural inputs, eligibility for social benefits—at speeds and scales that exceed human capacity to audit or contest in real time. In wealthy democracies with robust regulatory institutions and active civil societies, this opacity is a serious concern. In countries with weak rule of law, limited institutional capacity, and constrained civic space, it is potentially catastrophic.

An AI credit system that denies a loan carries no obligation to explain its reasoning. An AI diagnostic tool that misclassifies a disease leaves no readily accessible paper trail for accountability or correction. An AI benefit-distribution system that systematically excludes particular ethnic groups—whether by deliberate design or by biased training data—may produce discriminatory outcomes invisible to any oversight mechanism that exists. The combination of AI's inherent opacity and developing-country governance deficits creates conditions for harms that are difficult to detect, impossible to contest through normal channels, and potentially irreversible in their effects on individuals and communities.

5 Toward Abundance by Design: A Framework

The critique is not a counsel for pessimism. It is a design specification. The question is not whether AI can reduce poverty—the evidence that it can is already substantial and growing. The question is what values, structures, and accountability mechanisms must be built into AI systems if they are to reduce poverty equitably, durably, and at the scale the problem demands.

Five principles emerge from the evidence.

1. Proximity of design. AI tools that effectively serve poor communities are, al-

most without exception, those designed *with* those communities rather than exported to them. The most successful precision agriculture tools have been co-created with smallholder farmers. The most effective health AI instruments have been trained on locally collected data and validated by community health workers who understand local disease ecology. Proximity is not sentimentality; it is a practical requirement for effective optimisation.

2. **Sovereignty over data.** The communities whose lives generate training data must hold legal and practical rights over that data, including the right to withhold it, to access models trained on it, and to share in revenues generated from it. Data sovereignty is not merely an ethical aspiration; it is a prerequisite for avoiding the extractive patterns that have characterised previous eras of development.

3. **Legibility of decisions.** Any AI system that determines people's access to resources—credit, healthcare, food, social benefits—must be required to furnish human-understandable explanations of those determinations, and must be subject to meaningful, accessible appeal mechanisms. This is simultaneously a technical constraint, a regulatory requirement, and a design imperative.

4. **Deliberate redistribution of productivity gains.** The economic value created by AI-driven productivity improvements will not automatically flow to poor people. Achieving that requires explicit redistribution mechanisms: progressive taxation of AI-generated profits, sovereign wealth funds capitalised by data-exporting nations, and community ownership stakes in shared AI infrastructure. Efficiency gains must be consciously converted into poverty reduction rather than allowed to deepen wealth concentration.

5. **Humility about what optimisation can achieve.** No system of algorithms can resolve the political dimensions of poverty. Land-tenure reform, public health system building, anti-corruption enforcement, infrastructure investment—these require political will, institutional capacity, and sustained long-term commitment that no codebase can substitute for. AI is a lever; it is not the fulcrum on which the history of poverty turns.

The goal is not a world administered by algorithms. It is a world in which human judgement is amplified, distributed, and democratised—so that the wisdom of the best agronomist, physician, or credit analyst becomes accessible to everyone.

6 Conclusion: Designing the Abundance We Choose

There is a version of AI-driven development that looks like previous eras of technological optimism: brilliant instruments applied to complex problems by well-intentioned outsiders, producing impressive aggregate statistics and cascading unintended consequences

that take decades to fully surface. Development history has produced this version before. There is no structural reason AI cannot produce it again.

There is also a version that looks different. In this version, AI is not imposed on poor communities from the outside but co-created with them. It is not used to make decisions *about* people but to place better information and richer options in the hands of people making decisions for themselves. It is not owned by distant corporations but governed by the communities it serves. It does not optimise for proxy metrics but for genuine outcomes—and it remains epistemically humble about the gap between them.

Amara Kone did not need an algorithm to tell her that her land mattered, or that her daughter deserved an education. She needed better information about when to plant, what her harvest was worth, and what price she could command in a market she had never previously been able to see. The AI gave her access to knowledge that wealthier farmers had always possessed. That is not disruption. That is justice.

The poverty of the Global South is not a natural feature of the landscape. It is the accumulated output of centuries of specific, nameable decisions—about trade law, about governance, about whose labour was valued and whose was not—decisions made by identifiable actors in identifiable places. AI did not make those decisions. But AI, designed with genuine intention and genuine accountability, has the potential to begin reversing their consequences at a speed and scale that no previous instrument has made possible.

The abundance is not inevitable. It is a design choice. And choices, unlike natural phenomena, can be made differently.

Sources and Key References

- World Bank Poverty and Inequality Platform (2024). *Poverty Headcount Ratio at $2.15 a Day.* Washington, D.C.: World Bank Group.
- PwC (2017). *Sizing the Prize: What's the Real Value of AI for Your Business and How Can You Capitalise?* London: PricewaterhouseCoopers.
- McKinsey Global Institute (2018). *Notes from the AI Frontier: Modeling the Impact of AI on the World Economy.* New York: McKinsey & Company.
- McKinsey Global Institute (2017). *Jobs Lost, Jobs Gained: Workforce Transitions in a Time of Automation.* New York: McKinsey & Company.
- Food and Agriculture Organization (2023). *The State of Food Security and Nutrition in the World.* Rome: FAO.
- World Health Organization (2023). *Global Health Observatory Data Repository.* Geneva: WHO.
- Liu, Y., et al. (2019). "Artificial intelligence-based breast cancer nodal metastasis detection." *Nature Medicine*, 25(3), 1133–1139.
- Microsoft AI for Good (2023). *AI for Agriculture: Programme Impact Report.* Redmond, WA: Microsoft Corporation.
- Farmerline (2023). *Agricultural Impact Assessment: Upper West Region Pilot.* Accra: Farmerline Ltd.

- GSMA Mobile Money Programme (2023). *State of the Industry Report on Mobile Money.* London: GSMA.
- Demirgüç-Kunt, A., et al. (2022). *The Global Findex Database 2021: Financial Inclusion, Digital Payments, and Resilience in the Age of COVID-19.* Washington, D.C.: World Bank Group.

The End