

The Complete Treatise on Agentic AI

Soumadeep Ghosh with Qwen3.7

Kolkata, India

Abstract

This treatise provides a comprehensive, multi-disciplinary analysis of Agentic Artificial Intelligence. By synthesizing mathematics, computer science, economics, philosophy, political science, and warfare economics, we establish a unified framework for understanding autonomous agents. We formalize agent decision-making via Markov Decision Processes, prove convergence properties of value iteration, analyze the macroeconomic impacts of algorithmic market makers, and model geopolitical AI arms races using modified Lanchester equations.

The paper ends with “The End”

1 Introduction and Philosophy of Agency

The transition from passive machine learning models to *Agentic AI* represents a paradigm shift in computational philosophy. An agent is not merely a function approximator; it is an entity exhibiting intentionality, autonomy, and goal-directed behavior.

Philosophically, this touches upon Searle’s Chinese Room argument. While a narrow AI merely manipulates symbols (syntax), an agentic AI operating in an open-ended environment must develop internal world models that map to semantic reality. We define an agent’s agency $\mathcal{A}$ as a function of its perceptual capacity P , action space $\mathcal{A}_{ct}$, and temporal horizon τ :

$$\mathcal{A} = \lim_{\tau \rightarrow \infty} \mathbb{E} \left[\sum_{t=0}^{\tau} \gamma^t R(s_t, a_t) \right]$$

where $\gamma \in (0, 1)$ is the discount factor representing the agent’s temporal preference, a concept bridging computer science and economic theory.

2 Mathematical Foundations and Algorithms

The behavior of Agentic AI is rigorously modeled using Markov Decision Processes (MDPs). An MDP is defined as a tuple $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, P, R, \gamma \rangle$, where $\mathcal{S}$ is the state space, $\mathcal{A}$ is the action space, $P(s'|s, a)$ is the transition kernel, and $R(s, a)$ is the reward function.

The optimal value function $V^*(s)$ satisfies the Bellman optimality equation:

$$V^*(s) = \max_{a \in \mathcal{A}} \left[R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s') \right]$$

2.1 Algorithmic Implementation: Proximal Policy Optimization

To solve high-dimensional MDPs, we employ deep reinforcement learning. Below is the pseudo-code for Proximal Policy Optimization (PPO), a standard algorithm for training stable agentic policies.

Algorithm 1 Agentic Policy Optimization (PPO)

Require: Initial policy parameters θ , value function parameters ϕ , clipping parameter ϵ , number of iterations N , number of epochs K

Ensure: Optimized policy parameters θ^*

```
1: for iteration = 1 to  $N$  do
2:   Run policy  $\pi_\theta$  in environment to collect trajectory data  $\mathcal{D}$ 
3:   Compute advantage estimates  $\hat{A}_t$  using Generalized Advantage Estimation (GAE)
4:   for epoch = 1 to  $K$  do
5:     for each minibatch in  $\mathcal{D}$  do
6:       Calculate probability ratio  $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ 
7:       Calculate clipped surrogate objective  $L^{CLIP}(\theta)$ 
8:       Update  $\theta$  via stochastic gradient ascent on  $L^{CLIP}$ 
9:       Update  $\phi$  via mean squared error regression on value targets
10:    end for
11:  end for
12: end for
```

3 Economics, Finance, and Data Science

In financial markets, Agentic AI acts as an autonomous market maker and algorithmic trader. The agent's objective is to maximize expected utility over a consumption stream c_t :

$$\max_{\{c_t\}} \mathbb{E} \left[\sum_{t=0}^{\infty} \beta^t u(c_t) \right]$$

where $u(c_t) = \frac{c_t^{1-\rho}}{1-\rho}$ is the CRRA utility function.

From a data science perspective, agents process high-dimensional alternative data (satellite imagery, NLP on financial transcripts) to update their beliefs. Using Bayesian inference, the posterior distribution of market states θ given data D is:

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{\int P(D|\theta')P(\theta')d\theta'}$$

3.1 Macroeconomic Impact of Agentic AI

Table 1 summarizes the projected macroeconomic shifts driven by autonomous agents.

Table 1: Macroeconomic Impacts of Agentic AI Deployment

Sector	Traditional Model	Agentic AI Model
Labor Market	Human-centric, wage rigidity	Task-based, dynamic reallocation
Market Efficiency	Bounded rationality, friction	Hyper-rational, zero-latency arbitrage
Capital Allocation	Heuristic, historical bias	Predictive, real-time optimization
Wealth Distribution	Labor-capital split	Capital-skewed, high Gini coefficient

4 Geopolitics, International Relations, and Warfare

The deployment of Agentic AI has triggered a new security dilemma in international relations. States are engaged in an AI arms race, modeled here as a non-zero-sum game.

4.1 Warfare Economics and Lanchester Equations

In modern warfare, the integration of autonomous drone swarms alters the classical Lanchester square law. Let $x(t)$ and $y(t)$ be the combat strengths of two opposing forces. The modified differential equations incorporating AI force multipliers (α, β) and autonomous attrition rates (a, b) are:

$$\begin{aligned}\frac{dx}{dt} &= -ay + \alpha_{AI}x \log(x) \\ \frac{dy}{dt} &= -bx + \beta_{AI}y \log(y)\end{aligned}$$

The logarithmic terms represent the network effects and exponential scaling of AI-driven swarm intelligence.

4.2 Geopolitical AI Alliances

Figure 1 illustrates the current geopolitical network of AI research alliances and technological rivalries.




Figure 1: Geopolitical Network of AI Alliances and Rivalries (2026). Solid green lines indicate technological alliances; dashed red lines indicate strategic rivalries; solid blue lines indicate regulatory/trade cooperation.

5 Statistical Proofs and Convergence

To guarantee the reliability of Agentic AI, we must prove the convergence of its underlying learning algorithms. We provide a formal proof of the contraction mapping property for the Bellman operator.

Theorem: *The Bellman optimality operator T is a γ -contraction in the sup-norm $\|\cdot\|_\infty$.*

Proof: Let $V, U \in \mathbb{R}^{|S|}$ be two value functions. The Bellman operator is defined as:

$$(TV)(s) = \max_a \left[R(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s') \right]$$

We evaluate the difference between $(TV)(s)$ and $(TU)(s)$:

$$|(TV)(s) - (TU)(s)| = \left| \max_a \left[R + \gamma \sum P V \right] - \max_a \left[R + \gamma \sum P U \right] \right|$$

Using the property that $|\max x - \max y| \leq \max |x - y|$, we obtain:

$$|(TV)(s) - (TU)(s)| \leq \max_a \left| \gamma \sum_{s'} P(s'|s, a) [V(s') - U(s')] \right|$$

Since $P(s'|s, a)$ is a probability distribution (sums to 1) and $\gamma > 0$:

$$|(TV)(s) - (TU)(s)| \leq \gamma \sum_{s'} P(s'|s, a) |V(s') - U(s')|$$

$$|(TV)(s) - (TU)(s)| \leq \gamma \sum_{s'} P(s'|s, a) \|V - U\|_\infty$$

$$|(TV)(s) - (TU)(s)| \leq \gamma \|V - U\|_\infty$$

Taking the supremum over all states $s \in \mathcal{S}$:

$$\|TV - TU\|_\infty \leq \gamma \|V - U\|_\infty$$

Since $\gamma \in (0, 1)$, T is a contraction mapping. By the Banach Fixed-Point Theorem, T has a unique fixed point V^* , and iterative application $V_{k+1} = TV_k$ converges to V^* . ■

6 Conclusion

This treatise demonstrates that autonomous agents are not merely computational tools, but complex entities that reshape economic markets, alter geopolitical balances, and redefine the philosophy of mind. As we advance, the mathematical guarantees of convergence and the ethical frameworks of alignment must remain at the forefront of computer science and public policy.

References

- [1] Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). MIT Press.
- [2] Russell, S. (2019). *Human Compatible: Artificial Intelligence and the Problem of Control*. Viking.
- [3] Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- [4] Kissinger, H., Schmidt, E., & Huttenlocher, D. (2021). *The Age of AI: And Our Human Future*. Little, Brown and Company.
- [5] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.
- [6] Acemoglu, D., & Restrepo, P. (2018). The race between man and machine: Implications of technology for growth, factor shares, and employment. *American Economic Review*, 108(6), 1488-1542.
- [7] Allison, G. (2017). *Destined for War: Can America and China Escape Thucydides's Trap?* Houghton Mifflin Harcourt.

Glossary

Agentic AI

Artificial intelligence systems capable of perceiving their environment, making autonomous decisions, and taking actions to achieve specific goals without continuous human intervention.

Markov Decision Process (MDP)

A mathematical framework for modeling decision making in situations where outcomes are partly random and partly under the control of a decision maker.

Bellman Equation

A necessary condition for optimality associated with the dynamic programming algorithm, which breaks down a complex optimization problem into simpler sub-problems.

Nash Equilibrium

A concept in game theory where each player's strategy is optimal, given the strategies of all other players; no player has an incentive to unilaterally change their action.

Proximal Policy Optimization (PPO)

A family of policy-gradient methods for reinforcement learning which strike a balance between ease of implementation, sample complexity, and ease of tuning.

Lanchester Equations

Differential equations describing the time dependence of two adversaries' strengths as a function of time, used here in a modified form to model AI-driven swarm warfare.

Bayesian Inference

A method of statistical inference in which Bayes' theorem is used to update the probability for a hypothesis as more evidence or information becomes available.

The End