

The Modern Secretary Problem

Lord Soumadeep Ghosh with ChatGPT (OpenAI)

Kolkata, India

Abstract

The classical Secretary Problem is a canonical model of sequential decision-making under uncertainty: candidates arrive in random order, rejected candidates cannot be recalled, and the decision-maker seeks to select the best candidate. Its celebrated asymptotic solution rejects approximately the first $1/e$ fraction and then selects the first subsequent record candidate, achieving success probability approaching $1/e$ [24, 14]. This paper develops a compact modern extension in which candidates possess multiple observable attributes, including education, experience, assessment scores, references, and physical appearance, while ultimate suitability is treated as a latent variable. The resulting problem is formulated as Bayesian optimal stopping with endogenous information acquisition. Connections are drawn to economics, finance, econometrics, statistics, psychology, biology, philosophy, political science, geopolitics, diplomacy, international relations, welfare economics, warfare economics, computer science, algorithms, and artificial intelligence. Particular attention is paid to the distinction between information and causal relevance, the value of information, exploration versus exploitation, fairness, institutional design, and the dangers of encoding social biases into automated selection systems.

The paper ends with “The End”

1 Introduction

The Secretary Problem asks how an agent should make an irrevocable choice when alternatives arrive sequentially and the quality of an alternative can only be evaluated relative to previous observations. The classical formulation is associated with the “marriage” or “best-choice” problem and has generated a substantial literature in probability, statistics, operations research, economics, and computer science [24, 8, 13].

The classical problem is deliberately austere. Candidates are effectively represented by a single latent ranking. A modern selection environment is considerably richer. A candidate may have education, professional experience, test scores, references, interpersonal characteristics, compensation requirements, and many other signals. Some signals may be causally relevant, some merely predictive, and others may influence decision-makers without being causally related to performance.

The modern problem is therefore not simply

Who is the best candidate? (1)

but rather

When should an agent stop acquiring information and make an irreversible choice? (2)

This formulation connects the Secretary Problem to dynamic programming, Bayesian decision theory, statistical learning, real options, sequential experimentation, and artificial intelligence [7, 11, 29, 19].

2 The Classical Secretary Problem

Let there be N candidates arriving in uniformly random order. Let the decision-maker observe candidates sequentially and either accept or reject each candidate immediately. Rejected candidates cannot subsequently be recalled.

The objective is

$$\max_{\tau} \Pr \left(Q_{\tau} = \max_{1 \leq i \leq N} Q_i \right), \quad (3)$$

where τ is the stopping time and Q_i denotes candidate quality.

Suppose the first r candidates are automatically rejected. Thereafter, the first candidate exceeding the maximum quality observed during the initial sample is selected. The success probability is

$$P_N(r) = \frac{r}{N} \sum_{k=r}^{N-1} \frac{1}{k}. \quad (4)$$

For large N ,

$$P_N(r) \approx \frac{r}{N} \ln \left(\frac{N}{r} \right). \quad (5)$$

Writing $x = r/N$ gives

$$P(x) = -x \ln x. \quad (6)$$

Hence

$$P'(x) = -(\ln x + 1), \quad (7)$$

and the maximum occurs at

$$x = \frac{1}{e}. \quad (8)$$

Thus

$$r^* \approx \frac{N}{e}, \quad P(r^*) \approx \frac{1}{e}. \quad (9)$$

The classical rule is consequently the celebrated “37% rule”.

3 From One Attribute to Many

A modern candidate can be represented by an observable vector

$$X_i = (Q_i, E_i, A_i, W_i, I_i, P_i, R_i), \quad (10)$$

where Q_i is latent suitability, E_i education, A_i physical attractiveness, W_i work experience, I_i an assessment score, P_i personality or interpersonal information, and R_i references.

The distinction between Q_i and the remaining variables is fundamental. The latter are signals concerning the former rather than necessarily being components of it.

A simple latent-variable representation is

$$Q_i = \beta_E E_i + \beta_W W_i + \beta_I I_i + \beta_P P_i + \varepsilon_i. \quad (11)$$

However, an observable attribute should not automatically be interpreted as a causal determinant. In particular, an appearance variable may affect a decision-maker's perception without affecting actual productivity:

$$\widehat{Q}_i = Q_i + \gamma A_i + \eta_i. \quad (12)$$

Here $\widehat{Q}_i$ represents perceived suitability rather than actual suitability. This distinction is important for econometrics, psychology, organizational economics, and algorithmic fairness [3, 6, 17, 4].

4 Bayesian Updating

Let the information available after observing candidate i be

$$\mathcal{F}_i = \sigma(X_1, \dots, X_i). \quad (13)$$

The posterior distribution of latent quality is

$$p(Q_i \mid X_i, \mathcal{F}_{i-1}). \quad (14)$$

The posterior expected quality is therefore

$$\mu_i = \mathbb{E}[Q_i \mid X_i, \mathcal{F}_{i-1}]. \quad (15)$$

The posterior variance is

$$\sigma_i^2 = \text{Var}(Q_i \mid X_i, \mathcal{F}_{i-1}). \quad (16)$$

Consequently, two candidates with identical posterior means need not be equivalent: uncertainty itself can matter whenever decisions are risk-sensitive [33, 11, 32].

5 The Modern Optimal-Stopping Problem

Let U_i denote the utility associated with selecting candidate i . A general objective is

$$\max_{\tau} \mathbb{E}[U_{\tau}]. \quad (17)$$

Define $V_t(\mathcal{F}_t)$ as the maximum expected value obtainable from period t onward. Acceptance gives

$$V_t^A = \mathbb{E}[U_t \mid \mathcal{F}_t], \quad (18)$$

while rejection gives

$$V_t^R = -c + \mathbb{E}[V_{t+1}(\mathcal{F}_{t+1}) \mid \mathcal{F}_t], \quad (19)$$

where c is the cost of continuing the search.

The Bellman equation is therefore

$$V_t(\mathcal{F}_t) = \max \{V_t^A, V_t^R\}. \quad (20)$$

The optimal policy is

$$D_t^* = \begin{cases} 1, & V_t^A \geq V_t^R, \\ 0, & V_t^A < V_t^R. \end{cases} \quad (21)$$

The classical $1/e$ rule is therefore a special case of a substantially more general dynamic-programming problem [7, 29, 8].

6 Pareto Structure

With multiple attributes there need not be a total ordering. Candidate i Pareto dominates candidate j if

$$X_{ik} \geq X_{jk} \quad \text{for all relevant } k, \quad (22)$$

with strict inequality for at least one attribute.

A candidate with strong experience and weak test performance may therefore be incomparable with a candidate having weak experience and exceptional test performance.

The resulting feasible set has a Pareto frontier rather than a single scalar ranking [28, 10].

7 Information Versus Causation

The modern Secretary Problem requires a distinction between three concepts:

$$\text{predictive information} \neq \text{causal effect} \neq \text{decision-maker perception}. \quad (23)$$

A variable can improve prediction without being causally responsible for performance. Conversely, a variable can causally affect performance but be poorly measured.

A structural causal representation can be illustrated as follows.




Figure 1: Structural causal representation.

This distinction prevents an econometric correlation from being automatically interpreted as a structural production function.

8 Econometric Interpretation

Suppose realized performance is observed after hiring. An econometrician may estimate

$$Y_i = \beta_0 + \beta_E E_i + \beta_W W_i + \beta_I I_i + \beta_P P_i + \beta_A A_i + u_i. \quad (24)$$

However, observational estimates may suffer from omitted variables, selection bias, measurement error, simultaneity, and endogenous matching.

The Secretary Problem creates an additional selection mechanism because

$$\Pr(D_i = 1 \mid X_i, \mathcal{F}_i) \neq \Pr(D_i = 1). \quad (25)$$

Consequently, the observed hired population is not generally a random sample of applicants.

This connects sequential selection to the broader econometric literature on sample selection, treatment effects, and causal inference [18, 31, 1].

9 Psychology and Behavioral Decision-Making

The normative optimum need not describe actual human behavior. Human decision-makers may display anchoring, availability effects, representativeness heuristics, loss aversion, overconfidence, or premature stopping [22, 21].

The decision-maker can therefore be represented as having a subjective posterior

$$\tilde{p}(Q_i | X_i, \mathcal{F}_i) \neq p(Q_i | X_i, \mathcal{F}_i). \quad (26)$$

The resulting behavioral stopping rule is

$$\tilde{V}_t^A \geq \tilde{V}_t^R, \quad (27)$$

rather than the normative Bayesian rule.

Psychiatric concepts should be used cautiously here: deviations from an idealized rational policy do not by themselves constitute psychiatric disorders. Clinical diagnosis requires substantially different evidence and criteria. The relevant scientific domain here is primarily behavioral economics and cognitive psychology.

10 Biology and Evolution

Sequential selection also has biological analogues. Organisms search for mates, territories, food, and resources under uncertainty and finite time. Optimal foraging theory studies related exploration and exploitation trade-offs [25, 38].

A simplified biological objective can be represented as

$$\max_{\tau} \mathbb{E}[B(X_{\tau}) - C(\tau)], \quad (28)$$

where B denotes biological benefit and C search cost.

The Secretary Problem can consequently be interpreted as a stylized mathematical model of sequential ecological choice.

11 Physics: Search as a Stochastic Process

The process can be interpreted as a finite-horizon stochastic system with state

$$S_t = (t, \mathcal{F}_t). \quad (29)$$

Transitions are generated by the arrival of a new candidate:

$$S_t \longrightarrow S_{t+1}. \quad (30)$$

The analogy with statistical mechanics is limited but useful: the system evolves through a state space, information changes the state distribution, and the stopping rule selects an absorbing state.

The analogy should not be confused with a physical law. The Secretary Problem is a decision model, not a thermodynamic theory.

12 Chemistry and Materials Selection

The same mathematical structure appears in sequential screening of molecules, catalysts, materials, or experimental formulations.

Let Z_i be the experimentally observed property vector of compound i . The objective is

$$\max_{\tau} \mathbb{E}[U(Z_{\tau})]. \quad (31)$$

Experimental measurement may itself be costly, creating

$$\text{expected improvement} > \text{measurement cost} \quad (32)$$

as a stopping criterion. This connects the Secretary Problem to sequential experimental design and Bayesian optimization [20, 37].

13 Finance and Real Options

Financial decision-making provides a particularly natural application. An investor may observe investment opportunities sequentially while facing an expiration date.

Let V_t denote the value of an opportunity and C_t the value of continuing to search. The exercise rule is

$$\text{Exercise at } t \iff V_t \geq C_t. \quad (33)$$

This is structurally analogous to an American option: exercising terminates the opportunity to wait, while waiting preserves the option to obtain future information [27, 12].

The Secretary Problem can therefore be interpreted as a particularly stark real-option problem in which the exercise decision is irreversible.

14 Welfare Economics and Social Choice

Individual optimal stopping need not maximize social welfare.

Suppose candidate i produces social surplus

$$W_i = U_i^{\text{decision-maker}} + U_i^{\text{workers}} + U_i^{\text{consumers}} + U_i^{\text{externalities}}. \quad (34)$$

Then private and social stopping rules can differ:

$$\arg \max_{\tau} \mathbb{E}[U_{\tau}] \neq \arg \max_{\tau} \mathbb{E}[W_{\tau}]. \quad (35)$$

This distinction is central to welfare economics and mechanism design [2, 35, 26].

15 Political Science, Diplomacy, and International Relations

The same structure appears in repeated policy and diplomatic decisions. A state may receive successive offers, coalition proposals, treaties, or strategic signals.

Let

$$X_t = (\text{military, economic, diplomatic, domestic, informational}) \quad (36)$$

represent the state observed at time t .

A decision-maker must compare immediate strategic value with the continuation value of waiting.

The appropriate interpretation is institutional rather than normative: the framework describes a sequential decision problem and does not determine which political objective a government ought to pursue.

16 Warfare Economics and Strategic Competition

Military decision-making also contains sequential stopping structures: reconnaissance, target selection, procurement, mobilization, alliance formation, and resource allocation can involve costly information acquisition.

A generic objective is

$$\max_{\tau} \mathbb{E}[B_{\tau} - C_{\tau} - R_{\tau}], \quad (37)$$

where B_{τ} represents benefit, C_{τ} direct resource cost, and R_{τ} risk.

Unlike the classical Secretary Problem, strategic environments involve interacting agents. Consequently, the probability distribution of future opportunities may itself depend on the decisions already made.

This transforms the single-agent problem into a stochastic game [41, 15].

17 Information Acquisition

Suppose the decision-maker can purchase an additional signal Z_i at cost c_Z .

The value of information is

$$VOI = V^{\text{with information}} - V^{\text{without information}}. \quad (38)$$

Information acquisition is optimal when

$$VOI > c_Z. \quad (39)$$

Thus the modern Secretary Problem contains two stopping decisions:

1. whether to continue searching for another candidate;
2. whether to acquire additional information about the current candidate.

This produces a general sequential decision architecture rather than a fixed rejection fraction.

18 Machine Learning

A machine-learning implementation can estimate

$$\hat{\mu}_i = f_{\theta}(X_i, \mathcal{F}_{i-1}), \quad (40)$$

where f_{θ} may be a regression model, tree ensemble, neural network, Gaussian process, or another probabilistic learner.

The stopping policy can then be

$$D_i = \mathbb{I}[\hat{\mu}_i \geq \hat{C}_i]. \quad (41)$$

The system should preferably output uncertainty as well as a point estimate:

$$(\hat{\mu}_i, \hat{\sigma}_i^2). \quad (42)$$

This makes Bayesian neural networks, Gaussian processes, conformal prediction, and distributional learning relevant extensions [30, 5, 40].

19 Algorithmic Formulation

A generic finite-horizon algorithm is:

1. Observe the candidate vector X_t .
2. Update the posterior distribution of latent quality.
3. Estimate expected utility.
4. Estimate continuation value.
5. Compare acceptance and continuation values.
6. Accept if the acceptance value is at least the continuation value.
7. Otherwise reject and continue.

A compact pseudocode representation is:

Input: candidate stream $X[1], \dots, X[N]$

Initialize posterior belief B_0

Initialize continuation model C

for $t = 1, \dots, N$:

 observe $X[t]$

 update B_t using $X[t]$

 estimate current utility U_t

 estimate continuation value C_t

 if $U_t \geq C_t$:

 accept candidate t

 stop

 otherwise:

 reject candidate t

if no candidate was accepted:

 apply terminal decision rule

The algorithm is deliberately abstract because the statistical model and objective function determine the appropriate posterior and continuation-value estimators.

20 Artificial Intelligence and Reinforcement Learning

The problem can be represented as a Markov decision process

$$\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, R, \gamma), \tag{43}$$

where the state contains the current candidate, historical observations, posterior beliefs, and remaining horizon.

The optimal value function satisfies

$$V(s) = \max_{a \in \mathcal{A}} [R(s, a) + \gamma \mathbb{E} [V(s') \mid s, a]] . \quad (44)$$

For finite-horizon selection, γ may be replaced by an explicitly time-dependent horizon.

Modern reinforcement learning therefore provides computational tools for approximating the optimal policy when an analytical solution is unavailable [39].

21 Exploration and Exploitation

The classical $1/e$ strategy already contains an exploration-exploitation trade-off.

The modern formulation makes it explicit:

$$\text{exploration} \longleftrightarrow \text{exploitation}. \quad (45)$$

Exploration generates information about the candidate distribution. Exploitation converts current information into an immediate decision.

If model parameters are uncertain, continued observation can have informational value even when the next candidate is not itself attractive.

This connects the Secretary Problem to multi-armed bandits and Bayesian experimental design [16, 34].

22 Fairness and Institutional Design

A selection algorithm may be statistically predictive while still producing undesirable social outcomes. A model must distinguish predictive performance from normative fairness.

Possible constraints include demographic parity, equalized error rates, calibration, individual fairness, or explicitly welfare-based objectives. These criteria can conflict with one another under realistic statistical conditions [23, 9].

The relevant optimization problem can therefore be written as

$$\max_{\pi} \mathbb{E}[W(\pi)] \quad (46)$$

subject to

$$G_j(\pi) \leq b_j, \quad j = 1, \dots, m, \quad (47)$$

where π is the stopping policy and G_j represents institutional constraints.

This makes the modern Secretary Problem simultaneously a statistical and governance problem.

23 Philosophical Interpretation

The problem illustrates several classical philosophical distinctions.

First,

$$\text{epistemic uncertainty} \neq \text{ontological uncertainty}. \quad (48)$$

The decision-maker may simply lack information about a candidate whose quality is already fixed.

Second,

$$\text{prediction} \neq \text{explanation}. \quad (49)$$

A variable can improve prediction without explaining the underlying mechanism.
Third,

$$\text{rational decision} \neq \text{certain knowledge.} \quad (50)$$

Optimal stopping is precisely the theory of action when knowledge is incomplete [33, 36].

24 Religion and Ethical Decision-Making

Religious traditions provide different philosophical perspectives on uncertainty, prudence, duty, and the consequences of action. Such perspectives can be represented as alternative utility or constraint structures rather than as empirical parameters.

For example, a deontological formulation emphasizes admissible actions,

$$a \in \mathcal{A}_{\text{permissible}}, \quad (51)$$

whereas a consequentialist formulation emphasizes expected outcomes,

$$a^* = \arg \max_a \mathbb{E}[U \mid a]. \quad (52)$$

Virtue-oriented approaches instead evaluate characteristics of the decision-maker and the decision process.

The mathematical framework can therefore represent different ethical objectives without claiming that mathematics alone resolves ethical questions.

25 A Unified Conceptual Framework

The interdisciplinary structure can be summarized as follows.




Figure 2: Interdisciplinary structure.

The modern Secretary Problem is consequently a feedback system:

$$\text{Observation} \rightarrow \text{Belief} \rightarrow \text{Decision} \rightarrow \text{Outcome} \rightarrow \text{Learning.} \quad (53)$$

26 Comparison of Models

Table 1: Evolution of the Secretary Problem

Model	Information	Decision rule
Classical Secretary Problem	Relative one-dimensional rank	Reject approximately N/e , then accept first record
Multivariate problem	Several observable attributes	Attribute-dependent acceptance region
Bayesian problem	Signals about latent quality	Compare posterior value with continuation value
Information-acquisition problem	Candidate data plus purchasable signals	Acquire information when its value exceeds its cost
Machine-learning problem	Learned predictive model	Compare predicted value with estimated continuation value
Strategic problem	Endogenous multi-agent environment	Dynamic-game stopping policy
Social-selection problem	Utility plus institutional constraints	Constrained welfare-maximizing policy

27 Conclusion

The classical Secretary Problem is not merely a puzzle about interviewing candidates. It is a fundamental model of irreversible sequential choice.

Its classical solution,

$$r^* \approx \frac{N}{e}, \quad (54)$$

is an elegant special case. Once candidates possess heterogeneous signals and the decision-maker can learn about latent quality, the natural object becomes a Bayesian dynamic program:

$$V_t = \max \{V_t^A, V_t^R\}. \quad (55)$$

The central economic quantity is therefore not simply candidate quality, but the **option value of waiting for additional information**.

The modern formulation also exposes an important methodological principle. Education, experience, assessment scores, appearance, references, or any other observable variable should not automatically be interpreted as causal determinants of productivity. Statistical prediction, causal explanation, perception, and social evaluation are distinct objects.

Consequently, the Modern Secretary Problem is best understood as a general framework for

1. sequential choice under uncertainty
2. information acquisition
3. irreversibility, and
4. competing objectives

This framework applies across markets, organizations, science, technology, biology, strategic interaction, public policy, and artificial intelligence. Its deepest lesson is that rational choice under uncertainty is fundamentally a problem of deciding not merely *what* to choose, but *when enough information has been acquired to justify stopping*.

References

- [1] J. D. Angrist and J.-S. Pischke, *Mostly Harmless Econometrics*, Princeton University Press, 2009.
- [2] K. J. Arrow, *Social Choice and Individual Values*, Wiley, 1951.
- [3] K. J. Arrow, The theory of discrimination, in O. Ashenfelter and A. Rees, eds., *Discrimination in Labor Markets*, Princeton University Press, 1973.
- [4] S. Barocas, M. Hardt and A. Narayanan, *Fairness and Machine Learning*, fairmlbook.org, 2019.
- [5] D. Barber, *Bayesian Reasoning and Machine Learning*, Cambridge University Press, 2012.
- [6] G. S. Becker, *The Economics of Discrimination*, University of Chicago Press, 1957.
- [7] R. Bellman, *Dynamic Programming*, Princeton University Press, 1957.
- [8] F. T. Bruss, Sum the odds to one and stop, *Annals of Probability*, 28(3), 2000, 1384–1391.
- [9] A. Chouldechova, Fair prediction with disparate impact: A study of bias in recidivism prediction instruments, *Big Data*, 5(2), 2017, 153–163.
- [10] K. Deb, *Multi-Objective Optimization Using Evolutionary Algorithms*, Wiley, 2001.
- [11] M. H. DeGroot, *Optimal Statistical Decisions*, McGraw-Hill, 1970.
- [12] A. K. Dixit and R. S. Pindyck, *Investment Under Uncertainty*, Princeton University Press, 1994.
- [13] T. S. Ferguson, Who solved the secretary problem?, *Statistical Science*, 4(3), 1989, 282–296.
- [14] P. R. Freeman, The secretary problem and its extensions: A review, *International Statistical Review*, 51(2), 1983, 189–206.
- [15] D. Fudenberg and J. Tirole, *Game Theory*, MIT Press, 1991.
- [16] J. C. Gittins, Bandit processes and dynamic allocation indices, *Journal of the Royal Statistical Society, Series B*, 41(2), 1979, 148–164.
- [17] C. Goldin and C. Rouse, Orchestrating impartiality: The impact of “blind” auditions on female musicians, *American Economic Review*, 90(4), 2000, 715–741.
- [18] J. J. Heckman, Sample selection bias as a specification error, *Econometrica*, 47(1), 1979, 153–161.
- [19] R. A. Howard, Information value theory, *IEEE Transactions on Systems Science and Cybernetics*, 2(1), 1966, 22–26.
- [20] D. R. Jones, M. Schonlau and W. J. Welch, Efficient global optimization of expensive black-box functions, *Journal of Global Optimization*, 13, 1998, 455–492.
- [21] D. Kahneman, *Thinking, Fast and Slow*, Farrar, Straus and Giroux, 2011.
- [22] D. Kahneman and A. Tversky, Prospect theory: An analysis of decision under risk, *Econometrica*, 47(2), 1979, 263–291.

- [23] J. Kleinberg, S. Mullainathan and M. Raghavan, Inherent trade-offs in the fair determination of risk scores, *Proceedings of the 8th Innovations in Theoretical Computer Science Conference*, 2017.
- [24] D. V. Lindley, Dynamic programming and decision theory, *Applied Statistics*, 10(1), 1961, 39–51.
- [25] R. H. MacArthur and E. R. Pianka, On optimal use of a patchy environment, *American Naturalist*, 100(916), 1966, 603–609.
- [26] A. Mas-Colell, M. D. Whinston and J. R. Green, *Microeconomic Theory*, Oxford University Press, 1995.
- [27] R. C. Merton, Theory of rational option pricing, *Bell Journal of Economics and Management Science*, 4(1), 1973, 141–183.
- [28] V. Pareto, *Cours d’economie politique*, F. Rouge, 1896.
- [29] M. L. Puterman, *Markov Decision Processes*, Wiley, 1994.
- [30] C. E. Rasmussen and C. K. I. Williams, *Gaussian Processes for Machine Learning*, MIT Press, 2006.
- [31] D. B. Rubin, Estimating causal effects of treatments in randomized and nonrandomized studies, *Journal of Educational Psychology*, 66(5), 1974, 688–701.
- [32] D. B. Rubin, *Multiple Imputation for Nonresponse in Surveys*, Wiley, 1987.
- [33] L. J. Savage, *The Foundations of Statistics*, Wiley, 1954.
- [34] S. L. Scott, A modern Bayesian look at the multi-armed bandit, *Applied Stochastic Models in Business and Industry*, 26(6), 2010, 639–658.
- [35] A. K. Sen, *Collective Choice and Social Welfare*, Holden-Day, 1970.
- [36] H. A. Simon, A behavioral model of rational choice, *Quarterly Journal of Economics*, 69(1), 1955, 99–118.
- [37] J. Snoek, H. Larochelle and R. P. Adams, Practical Bayesian optimization of machine learning algorithms, in *Advances in Neural Information Processing Systems*, 2012.
- [38] D. W. Stephens and J. R. Krebs, *Foraging Theory*, Princeton University Press, 1986.
- [39] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed., MIT Press, 2018.
- [40] V. Vovk, A. Gammerman and G. Shafer, *Algorithmic Learning in a Random World*, Springer, 2005.
- [41] J. von Neumann and O. Morgenstern, *Theory of Games and Economic Behavior*, Princeton University Press, 1944.

Glossary

Bayesian updating

Revision of probability beliefs after observing new evidence.

Continuation value

Expected value of rejecting the current option and retaining the option to make a future decision.

Decision surface

The boundary in a multidimensional attribute space separating acceptance from rejection.

Exploration

Deliberate acquisition of information before committing to an action.

Exploitation

Using currently available information to obtain an immediate payoff.

Latent variable

An unobserved variable inferred indirectly from observable signals.

Optimal stopping

Mathematical theory of determining when a sequential process should be terminated.

Pareto dominance

A state is Pareto dominant when it is at least as good in every relevant dimension and strictly better in at least one.

Posterior distribution

Probability distribution of an unknown quantity conditional on observed information.

Secretary Problem

Sequential selection problem in which alternatives arrive randomly, rejected alternatives cannot be recalled, and the decision-maker seeks an optimal stopping rule.

Stopping time

A random time whose occurrence can be determined using information available up to that time.

Value of information

Expected improvement in decision value generated by acquiring additional information.

The End